

Introduction to MATLAB Programming

KEY TERMS

computer program	comments	modes
scripts	block comment	writing to a file
live script	comment blocks	appending to a file
algorithm	input/output (I/O)	reading from a file
modular program	user	user-defined functions
top-down design	empty array	function call
external file	error message	argument
default input device	formatting	control
prompting	format specifier	return value
default output device	place holder	function header
execute/run	conversion characters	output arguments
high-level languages	newline character	input arguments
machine language	field width	function body
executable	leading blanks	function definition
compiler	trailing zeros	local variables
source code	plot symbols	scope of variables
object code	markers	base workspace
interpreter	line types	
documentation	toggle	

CONTENTS

3.1 Algorithms ...	80
3.2 MATLAB Scripts	81
3.3 Input and Output	85
3.4 Scripts With Input and Output	93
3.5 Scripts to Produce and Customize Simple Plots	94
3.6 Introduction to File Input/Output (Load and Save)	101
3.7 User-Defined Functions That Return a Single Value	106
3.8 Local Functions in Scripts	115
3.9 Commands and Functions ...	116
Summary	118

We have now used the MATLAB[®] product interactively in the Command Window. That is sufficient when all one needs is a simple calculation. However, in many cases, quite a few steps are required before the final result can be obtained. In those cases, it is more convenient to group statements together in what is called a *computer program*.

In this chapter, we will introduce the simplest MATLAB programs, which are called *scripts*. Examples of scripts that customize simple plots will illustrate

Common Pitfalls	118
Programming Style Guidelines	118

the concept. Input will be introduced, both from files and from the user. Output to files and to the screen will also be introduced. Finally, user-defined functions that calculate and return a single value will be described. These topics serve as an introduction to programming, which will be expanded on in [Chapter 6](#).

As of Version R2016a, there are two types of scripts. A new, richer script type called a *live script* has been created in MATLAB. Live scripts will be introduced in [Chapter 6](#), in which programming concepts will be covered in more depth. In this chapter, we will create simple scripts stored in MATLAB code files, which have an extension of .m.

3.1 ALGORITHMS

Before writing any computer program, it is useful to first outline the steps that will be necessary. An *algorithm* is the sequence of steps needed to solve a problem. In a *modular* approach to programming, the problem solution is broken down into separate steps, and then each step is further refined until the resulting steps are small enough to be manageable tasks. This is called the *top-down design* approach.

As a simple example, consider the problem of calculating the area of a circle. First, it is necessary to determine what information is needed to solve the problem, which in this case is the radius of the circle. Next, given the radius of the circle, the area of the circle would be calculated. Finally, once the area has been calculated, it has to be displayed in some way. The basic algorithm then is three steps:

- Get the input: the radius
- Calculate the result: the area
- Display the output

Even with an algorithm this simple, it is possible to further refine each of the steps. When a program is written to implement this algorithm, the steps would be as follows.

- Where does the input come from? Two possible choices would be from an *external file*, or from the user (the person who is running the program) who enters the number by typing it from the keyboard. For every system, one of these will be the *default input device* (which means, if not specified otherwise, this is where the input comes from!). If the user is supposed to enter the radius, the user has to be told to type in the radius (and, in what units). Telling the user what to enter is called *prompting*. So, the input step actually becomes two steps: prompt the user to enter a radius, and then read it into the program.

- To calculate the area, the formula is needed. In this case, the area of the circle is π multiplied by the square of the radius. So, that means the value of the constant for π is needed in the program.
- Where does the output go? Two possibilities are: (1) to an external file, or (2) to the screen. Depending on the system, one of these will be the *default output device*. When displaying the output from the program, it should always be as informative as possible. In other words, instead of just printing the area (just the number), it should be printed in a nice sentence format. Also, to make the output even more clear, the input should be printed. For example, the output might be the sentence: “For a circle with a radius of 1 inch, the area is 3.1416 inches squared.”

For most programs, the basic algorithm consists of the three steps that have been outlined:

1. Get the input(s)
2. Calculate the result(s)
3. Display the result(s)

As can be seen here, even the simplest problem solutions can then be refined further. This is top-down design.

3.2 MATLAB SCRIPTS

Once a problem has been analyzed, and the algorithm for its solution has been written and refined, the solution to the problem is then written in a particular programming language. A computer program is a sequence of instructions, in a given language, which accomplishes a task. To *execute* or *run* a program is to have the computer actually follow these instructions sequentially.

High-level languages have English-like commands and functions, such as “print this” or “if $x < 5$ do something.” The computer, however, can only interpret commands written in its *machine language*. Programs that are written in high-level languages must therefore be translated into machine language before the computer can actually execute the sequence of instructions in the program. A program that does this translation from a high-level language to an *executable* file is called a *compiler*. The original program is called the *source code*, and the resulting executable program is called the *object code*. Compilers translate from the source code to object code; this is then executed as a separate step.

By contrast, an *interpreter* goes through the code line-by-line, translating and executing each command as it goes. MATLAB uses what are called either script files or MATLAB code files, which have an extension on the file name of *.m*.

These script files are interpreted, rather than compiled. Therefore, the correct terminology is that these are scripts, not programs. However, the terms are used somewhat loosely by many people, and documentation in MATLAB itself refers to scripts as programs. In this book, we will reserve the use of the word “program” to mean a set of scripts and functions, as described briefly in [Section 3.7](#) and then in more detail in [Chapter 6](#).

A script is a sequence of MATLAB instructions that is stored in a file with an extension of .m and saved. The contents of a script can be displayed in the Command Window using the `type` command. The script can be executed, or run, by simply entering the name of the file (without the .m extension).

Before creating a script, make sure the Current Folder is set to the folder in which you want to save your files.

The steps involved in creating a script depend on the version of MATLAB. The easiest method is to click on “New Script” under the HOME tab. Alternatively, one can click on the down arrow under “New” and then choose Script (see [Figure 3.1](#))

A new window will appear called the Editor (which can be docked). In the latest versions of MATLAB, this window has three tabs: “EDITOR”, “PUBLISH”, and “VIEW”. Next, simply type the sequence of statements (note that line numbers will appear on the left).

When finished, save the file by choosing the Save down arrow under the EDITOR tab. Make sure that the extension of .m is on the file name (this should be the default). The rules for file names are the same as for variables (they must start with a letter; after that there can be letters, digits, or the underscore.)

If you have entered commands in the Command Window and decide that you would like to put them into a script, an alternate method for creating a

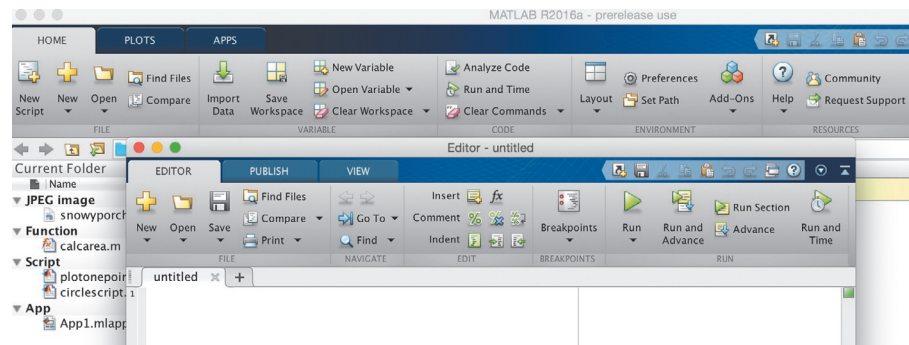


FIGURE 3.1
Toolstrip and Editor.

script is to select the commands in the Command History window, and then right click. This will give options for creating a script or live script and will then repopulate the editor with those commands.

In our first example, we will now create a script called *script1.m* that calculates the area of a circle. It assigns a value for the radius, and then calculates the area based on that radius.

In this book, scripts will be displayed in a box with the name of the file on top.

script1.m

```
radius = 5
area = pi * (radius ^2)
```

There are two ways to view a script once it has been written: either open the Editor Window to view it or use the **type** command, as shown here, to display it in the Command Window. The **type** command shows the contents of the file named *script1.m*; notice that the .m is not included:

```
>> type script1
radius = 5
area = pi * (radius ^2)
```

To actually run or execute the script from the Command Window, the name of the file is entered at the prompt (again, without the .m). When executed, the results of the two assignment statements are displayed, as the output was not suppressed for either statement.

```
>> script1
radius =
    5
area =
   78.5398
```

Once the script has been executed, you may find that you want to make changes to it (especially if there are errors!). To edit an existing file, there are several methods to open it. The easiest are:

- Within the Current Folder Window, double-click on the name of the file in the list of files.
- Choosing the Open down arrow will show a list of Recent Files

3.2.1 Documentation

It is very important that all scripts be *documented* well, so that people can understand what the script does and how it accomplishes its task. One way of documenting a script is to put *comments* in it. In MATLAB, a comment is anything from a % to the end of that particular line. Comments are completely ignored

when the script is executed. To put in a comment, simply type the % symbol at the beginning of a line, or select the comment lines and then click on the Edit down arrow and click on the % symbol, and the Editor will put in the % symbols at the beginning of those lines for the comments.

For example, the previous script to calculate the area of a circle could be modified to have comments:

```
circlescript.m

% This script calculates the area of a circle

% First the radius is assigned
radius = 5

% The area is calculated based on the radius
area = pi * (radius ^2)
```

The first comment at the beginning of the script describes what the script does; this is sometimes called a **block comment**. Then, throughout the script, comments describe different parts of the script (not usually a comment for every line, however!). Comments don't affect what a script does, so the output from this script would be the same as for the previous version.

The **help** command in MATLAB works with scripts as well as with built-in functions. The first block of comments (defined as contiguous lines at the beginning) will be displayed. For example, for *circlescript*:

```
>> help circlescript
This script calculates the area of a circle
```

The reason that a blank line was inserted in the script between the first two comments is that otherwise both would have been interpreted as one contiguous comment, and both lines would have been displayed with **help**. The very first comment line is called the “H1 line”; it is what the function **lookfor** searches through.

PRACTICE 3.1

Write a script to calculate the circumference of a circle ($C = 2\pi r$). Comment the script.

Longer comments, called **comment blocks**, consist of everything in between % { and % }, which must be alone on separate lines. For example:

```
%{
    this is used for a really
    Really
    REALLY
    long comment
%}
```

3.3 INPUT AND OUTPUT

The previous script would be much more useful if it were more general; for example, if the value of the radius could be read from an external source rather than being assigned in the script. Also, it would be better to have the script print the output in a nice, informative way. Statements that accomplish these tasks are called *input/output* statements, or *I/O* for short. Although, for simplicity, examples of input and output statements will be shown here in the Command Window, these statements will make the most sense in scripts.

3.3.1 Input Function

Input statements read in values from the default or standard input device. In most systems, the default input device is the keyboard, so the input statement reads in values that have been entered by the *user*, or the person who is running the script. To let the user know what he or she is supposed to enter, the script must first prompt the user for the specified values.

The simplest input function in MATLAB is called **input**. The **input** function is used in an assignment statement. To call it, a character vector is passed that is the prompt that will appear on the screen, and whatever the user types will be stored in the variable named on the left of the assignment statement. For ease of reading the prompt, it is useful to put a colon and then a space after the prompt. For example,

```
>> rad = input('Enter the radius: ')
Enter the radius: 5
rad =
    5
```

In this case, the prompt was printed and then the user entered 5.

If character or character vector input is desired, 's' must be added as a second argument to the **input** function:

```
>> letter = input('Enter a char: ', 's')
Enter a char: g
letter =
    'g'
```

If the user enters only spaces or tabs before hitting the Enter key, they are ignored and an *empty array* is stored in the variable:

```
>> mychar = input('Enter a character: ', 's')
Enter a character:
mychar =
    0×0 empty char array
```

However, if blank spaces are entered before other characters, they are included in the variable. In the next example, the user hit the space bar four times before entering “go.” The `length` function returns the number of characters in the variable.

```
>> mystr = input('Enter a word: ', 's')
Enter a word:      go
mystr =
      '      go'
>> length(mystr)
ans =
      6
```

QUICK QUESTION!

What would be the result if the user enters blank spaces after other characters? For example, the user here entered “xyz ” (four blank spaces):

Answer: The space characters would be stored in the variable.

```
>> mychar = input('Enter chars: ', 's')
Enter chars: xyz
mychar =
      'xyz      '
```

It is also possible for the user to type quotation marks around the entry rather than including the second argument ‘s’ in the call to the `input` function.

```
>> name = input('Enter your name: ')
Enter your name: 'Stormy'
name =
      'Stormy'
```

In this case, the user entered a character vector, but it is also possible to enter a string:

```
>> name = input('Enter your name: ')
Enter your name: "Stormy"
name =
      "Stormy"
```

However, this assumes that the user would know to do this, so it is better to signify that character input is desired in the `input` function itself. Also, if the ‘s’ is specified and the user enters quotation marks, these would become part of the variable.


```
>> name = input('Enter your name: ','s')
Enter your name: 'Stormy'
name =
    'Stormy'
>> length(name)
ans =
     8
```

Note what happens if character input has not been specified, but the user enters a letter rather than a number.

```
>> num = input('Enter a number: ')
Enter a number: t
Error using input
Unrecognized function or variable 't'.

Enter a number: 3
num =
     3
```

MATLAB gave an *error message* and repeated the prompt. However, if *t* is the name of a variable, MATLAB will take its value as the input.

```
>> t = 11;
>> num = input('Enter a number: ')
Enter a number: t
num =
    11
```

Separate **input** statements are necessary if more than one input is desired. For example,

```
>> x = input('Enter the x coordinate: ');
>> y = input('Enter the y coordinate: ');
```

Normally, in a script the results from **input** statements are suppressed with a semicolon at the end of the assignment statements.

PRACTICE 3.2

Create a script that would prompt the user for a length, and then 'f' for feet or 'm' for meters, and store both inputs in variables. For example, when executed, it would look like this (assuming the user enters 12.3 and then m):

```
Enter the length: 12.3
Is that f(eet) or m(eters)? : m
```

It is also possible to enter a vector. The user can enter any valid vector, using any valid syntax such as square brackets, the colon operator, or functions such as **linspace**.

```
>> v = input('Enter a vector: ')
Enter a vector: [3    8   22]
v =
     3     8    22
```

3.3.2 Output Statements: disp and fprintf

Output statements display character arrays or strings and/or the results of expressions and can allow for *formatting*, or customizing how they are displayed. The simplest output function in MATLAB is **disp**, which is used to display the result of an expression or a string or character array without assigning any value to the default variable *ans*. However, **disp** does not allow formatting. For example,

```
>> disp('Hello')
Hello

>> disp(4 ^ 3)
64
```

Formatted output can be printed to the screen using the **fprintf** function. For example,

```
>> fprintf('The value is %d, for sure!\n', 4 ^ 3)
The value is 64, for sure!

>>
```

To the **fprintf** function, first a string or character vector (called the *format specifier*) is passed that contains any text to be printed, as well as formatting information for the expressions to be printed. The format specifier can be either a string or a character vector; historically (before R2016b), it was always a character vector, so that is the way they will typically be shown, both in this book and in the MATLAB documentation. In this example, the %d is an example of format information.

Note: The format specifier in **fprintf** can be either a character vector or a string, but the prompt in an **input** statement must be a character vector.

The %d is sometimes called a *place holder* because it specifies where the value of the expression that is after the format specifier is to be printed. The character in the place holder is called the *conversion character*, and it specifies the type of value that is being printed. There are others, but what follows is a list of the simple place holders:

```
%d    integer (it stands for decimal integer)
%f    float (real number)
%c    character (one character)
%s    string of characters
```

Don't confuse the % in the place holder with the symbol used to designate a comment.

The character '\n' at the end of the format specifier is a special character called the *newline character*; what happens when it is printed is that the output that follows moves down to the next line.

QUICK QUESTION!

What do you think would happen if the newline character is omitted from the end of an **fprintf** statement?

Answer: Without it, the next prompt would end up on the same line as the output. It is still a prompt, and so an expression can be entered, but it looks messy as shown here.

```
>> fprintf('The value is %d, surely!', 4 ^3)
The value is 64, surely!>> 5 + 3
ans =
    8
```

Note that with the **disp** function, however, the prompt will always appear on the next line:

```
>> disp('Hi')
Hi
>>
```

Also, note that an ellipsis can be used after a string or character vector but not in the middle.

QUICK QUESTION!

How can you get a blank line in the output?

Answer: Have two newline characters in a row.

```
>> fprintf('The value is %d, \n\nOK!\n', 4 ^3)
The value is 64,
OK!
```

This also points out that the newline character can be anywhere in the format specifier; when it is printed, the output moves down to the next line.

Note that the newline character can also be used in the prompt in the **input** statement; for example:

```
>> x = input('Enter the \nx coordinate: ');
Enter the
x coordinate: 4
```

However, the newline is the **ONLY** formatting character allowed in the prompt in **input**.

To print two values, there would be two place holders in the format specifier, and two expressions after the format specifier. The expressions fill in for the place holders in sequence.

```
>> fprintf('The int is %d and the char is %c\n', ...
    33 - 2, 'x')
```

```
The int is 31 and the char is x
```

A **field width** can also be included in the place holder in **fprintf**, which specifies how many characters total are to be used in printing. For example, `%5d` would indicate a field width of 5 for printing an integer and `%10s` would indicate a field width of 10 for a string. For floats, the number of decimal places can also be specified; for example, `%6.2f` means a field width of 6 (including the decimal point and the two decimal places) with 2 decimal places. For floats, just the number of decimal places can also be specified; for example, `%.3f` indicates 3 decimal places, regardless of the field width.

```
>> fprintf('The int is %3d and the float is %6.2f\n', ...
    5, 4.9)
```

```
The int is 5 and the float is 4.90
```

Note that if the field width is wider than necessary, **leading blanks** are printed, and if more decimal places are specified than necessary, **trailing zeros** are printed.

QUICK QUESTION!

What do you think would happen if you tried to print 1234.5678 in a field width of 3 with 2 decimal places?

```
>> fprintf('%3.2f\n', 1234.5678)
```

Answer: It would print the entire 1234, but round the decimals to two places, that is,

```
1234.57
```

If the field width is not large enough to print the number, the field width will be increased. Basically, to cut the number off would give a misleading result, but rounding the decimal places does not change the number significantly.

QUICK QUESTION!

What would happen if you use the `%d` conversion character but you're trying to print a real number?

Answer: MATLAB will show the result using exponential notation

```
>> fprintf('%d\n', 1234567.89)
1.234568e+006
```

Note that if you want exponential notation, this is not the correct way to get it; instead, there are conversion characters that can be used. Use the **help** browser to see this option, as well as many others!

There are many other options for the format specifier. For example, the value being printed can be left-justified within the field width using a minus sign. The following example shows the difference between printing the integer 3 using %5d and using %-5d. The x's below are used to show the spacing.

```
>> fprintf('The integer is xx%5dxx and xx%-5dxx\n', 3, 3)
The integer is xx    3xx and xx3    xx
```

Also, strings and character vectors can be truncated by specifying “decimal places”:

```
>> fprintf('The string is %s or %.2s\n', "street", "street")
The string is street or st
```

There are several special characters that can be printed in the format specifier in addition to the newline character. To print a slash, two slashes in a row are used, and also to print a single quote two single quotes in a row are used. Additionally, ‘\t’ is the tab character.

```
>> fprintf('Try this out: tab\t quote \' \' slash \\\n')
Try this out: tab    quote ' slash \
```

3.3.2.1 Printing Vectors and Matrices

For a vector, if a conversion character and the newline character are in the format specifier, it will print in a column regardless of whether the vector itself is a row vector or a column vector.

```
>> vec = 2:5;
>> fprintf('%d\n', vec)
2
3
4
5
```

Without the newline character, it would print in a row, but the next prompt would appear on the same line:

```
>> fprintf('%d', vec)
2345>>
```

However, in a script, a separate newline character could be printed to avoid this problem. It is also much better to separate the numbers with spaces.

```
printvec.m
```

```
% This demonstrates printing a vector

vec = 2:5;
fprintf('%d ', vec)
fprintf('\n')
```

```
>> printvec
2 3 4 5
>>
```

If the number of elements in the vector is known, that many conversion characters can be specified and then the newline:

```
>> fprintf('%d %d %d %d\n', vec)
2 3 4 5
>>
```

This is not very general, however, and is therefore not preferable.

For matrices, MATLAB unwinds the matrix column by column. For example, consider the following 2×3 matrix:

```
>> mat = [5 9 8; 4 1 10]
mat =
     5     9     8
     4     1    10
```

Specifying one conversion character and then the newline character will print the elements from the matrix in one column. The first values printed are from the first column, then the second column, and so on.

```
>> fprintf('%d\n', mat)
5
4
9
1
8
10
```

If three of the %d conversion characters are specified, the **fprintf** will print three numbers across on each line of output, but again the matrix is unwound column-by-column. It again prints first the two numbers from the first column (across on the first row of output), then the first value from the second column, and so on.

```
>> fprintf('%d %d %d\n', mat)
5 4 9
1 8 10
```

If the transpose of the matrix is printed, however, using the three %d conversion characters, the matrix is printed as it appears when created.

```
>> fprintf('%d %d %d\n', mat') % Note the transpose
5 9 8
4 1 10
```

For vectors and matrices, even though formatting cannot be specified, the **disp** function may be easier to use in general than **fprintf** because it displays the result in a straightforward manner. For example,

```
>> mat = [15 11 14; 7 10 13]
mat =
    15    11    14
     7    10    13
>> disp(mat)
    15    11    14
     7    10    13
>> vec = 2:5
vec =
     2     3     4     5
>> disp(vec)
     2     3     4     5
```

Note that when loops are covered in [Chapter 5](#), formatting the output of matrices will be easier. For now, however, **disp** works well.

Note: there is another output function, **display**, which is generally not used by programmers and therefore will not be used in this text. The **display** function is used internally by MATLAB when an assignment statement is not suppressed. The **display** function is called implicitly in this case; it shows the variable name and the assignment operator, and generally, then calls the **disp** function to display the expression. The **disp** function does not show the variable name.

```
>> var = 33
var =
    33
>> display(var)
var =
    33
>> var
var =
    33
>> disp(var)
    33
```

When printing expressions, the output from **disp** and **display** may or may not be the same, depending on the type of the expression.

3.4 SCRIPTS WITH INPUT AND OUTPUT

Putting all of this together now, we can implement the algorithm from the beginning of this chapter. The following script calculates and prints the area

of a circle. It first prompts the user for a radius, reads in the radius, and then calculates and prints the area of the circle based on this radius.

circleIO.m

```
% This script calculates the area of a circle
% It prompts the user for the radius

% Prompt the user for the radius and calculate
% the area based on that radius
fprintf('Note: the units will be inches.\n')
radius = input('Please enter the radius: ');
area = pi * (radius ^2);

% Print all variables in a sentence format
fprintf('For a circle with a radius of %.2f inches,\n', ...
        radius)
fprintf('the area is %.2f inches squared\n', area)
```

Executing the script produces the following output:

```
>> circleIO
Note: the units will be inches.
Please enter the radius: 3.9
For a circle with a radius of 3.90 inches,
the area is 47.78 inches squared
```

Note that the output from the first two assignment statements (including the **input**) is suppressed by putting semicolons at the end. That is usually done in scripts, so that the exact format of what is displayed by the program is controlled by the **fprintf** functions.

PRACTICE 3.3

Write a script to prompt the user separately for a character and a number, and print the character in a field width of 3 and the number left justified in a field width of 8 with 3 decimal places. Test this by entering numbers with varying widths.

3.5 SCRIPTS TO PRODUCE AND CUSTOMIZE SIMPLE PLOTS

MATLAB has many graphing capabilities. Customizing plots is often desired and this is easiest to accomplish by creating a script rather than typing one command at a time in the Command Window. For that reason, simple plots and how to customize them will be introduced in this chapter on MATLAB programming.

The help topics that contain graph functions include **graph2d** and **graph3d**. Typing **help graph2d** would display some of the two-dimensional graph functions, as well as functions to manipulate the axes and to put labels and titles on the graphs. The Search Documentation under MATLAB Graphics also has a section on two- and three-dimensional plots.

3.5.1 The Plot Function

For now, we'll start with a very simple graph of one point using the **plot** function.

The following script, *plotonepoint*, plots one point. To do this, first values are given for the x and y coordinates of the point in separate variables. The point is plotted using a red star ('*'). The plot is then customized by specifying the minimum and maximum values on first the x and then y -axes. Labels are then put on the x -axis, the y -axis, and the graph itself using the functions **xlabel**, **ylabel**, and **title**. (Note: there are no default labels for the axes.)

All of this can be done from the Command Window, but it is much easier to use a script. The following shows the contents of the script *plotonepoint* that accomplishes this. The x coordinate represents the time of day (e.g., 11 a.m.) and the y coordinate represents the temperature (e.g., in degrees Fahrenheit) at that time.

plotonepoint.m

```
% This is a really simple plot of just one point!

% Create coordinate variables and plot a red '*'
x = 11;
y = 48;
plot(x,y,'r*')


% Change the axes and label them
axis([9 12 35 55])
xlabel('Time')
ylabel('Temperature')

% Put a title on the plot
title('Time and Temp')
```

In the call to the **axis** function, one vector is passed. The first two values are the minimum and maximum for the x -axis, and the last two are the minimum and maximum for the y -axis. Executing this script brings up a Figure Window with the plot (see [Figure 3.2](#)).

To be more general, the script could prompt the user for the time and temperature, rather than just assigning values. Then, the **axis** function could be used based on whatever the values of x and y are, as in the following example:

```
axis([x-2 x+2 y-10 y+10])
```

**FIGURE 3.2**

Plot of one data point.

In addition, although they are the x and y coordinates of a point, variables named *time* and *temp* might be more mnemonic than x and y .

PRACTICE 3.4


Modify the script *plotonepoint* to prompt the user for the time and temperature, and set the axes based on these values.

To plot more than one point, x and y vectors are created to store the values of the (x,y) points. For example, to plot the points

```
(1, 1)
(2, 5)
(3, 3)
(4, 9)
(5, 11)
(6, 8)
```

first an x vector is created that has the x values (as they range from 1 to 6 in steps of 1, the colon operator can be used) and then a y vector is created with the y values. The following will create (in the Command Window) x and y vectors and then plot them (see Figure 3.3).

```
>> x = 1:6;
>> y = [1 5 3 9 11 8];
>> plot(x,y)
```

**FIGURE 3.3**

Plot of data points from vectors.

Note that the points are plotted with straight lines drawn in between. Note also that the default color is a shade of blue. Also, the axes are set up according to the data; for example, the x values range from 1 to 6 and the y values from 1 to 11, so that is how the axes are set up. There are many options for the `axis` function; for example, just calling it with no arguments returns the values used for the x and y -axes ranges.

```
>> arang = axis
arang =
     1     6     1    11
```

Axes can also be turned on and off, and they can be made square or equal to each other. A subset of the data can also be shown by limiting the extent of the axes.

Also, note that in this case the x values are the indices of the y vector (the y vector has six values in it, so the indices iterate from 1 to 6). When this is the case, it is not necessary to create the x vector. For example,

```
>> plot(y)
```

will plot exactly the same figure without using an x vector.

3.5.1.1 Customizing a Plot: Color, Line Types, Marker Types

Plots can be done in the Command Window, as shown here, if they are really simple. However, many times it is desired to customize the plot with labels, titles, and so on, so it makes more sense to do this in a script. Using the `help`

function for **plot** will show the many options such as the line types and colors. In the previous script *plotonepoint*, the character vector 'r*' specified a red star for the point type. The *LineStyle*, or line specification, can specify up to three different properties in a character vector or string, including the color, line type, and the symbol or marker used for the data points.

The possible colors are:

```
b  blue
g  green
r  red
c  cyan
m  magenta
y  yellow
k  black
w  white
```

Either the single character listed above or the full name of the color can be used in the string to specify the color. The *plot symbols*, or *markers*, that can be used are:

```
.  point
o  circle
x  x-mark
+  plus
*  star
s  square
d  diamond
v  down triangle
^  up triangle
<  left triangle
>  right triangle
p  pentagram
h  hexagram
```

Line types can also be specified by the following:

```
-      solid
:      dotted
- .    dash dot
-      dashed
(none) no line
```

If no line type is specified and no marker type is specified, a solid line is drawn between the points, as seen in the last example.

3.5.2 Simple Related Plot Functions

Other functions that are useful in customizing plots include **clf**, **figure**, **hold**, **legend**, and **grid**. Brief descriptions of these functions are given here; use **help** to find out more about them:

clf: clears the Figure Window by removing everything from it.

figure: creates a new, empty Figure Window when called without any arguments. Calling it as **figure(n)** where *n* is an integer is a way of creating and maintaining multiple Figure Windows and of referring to each individually.

hold: is a toggle that freezes the current graph in the Figure Window, so that new plots will be superimposed on the current one. Just **hold** by itself is a *toggle*, so calling this function once turns the hold on, and then the next time turns it off. Alternatively, the commands **hold on** and **hold off** can be used.

legend: displays strings or character vectors passed to it in a legend box in the Figure Window, in order of the plots in the Figure Window

grid: displays grid lines on a graph. Called by itself, it is a toggle that turns the grid lines on and off. Alternatively, the commands **grid on** and **grid off** can be used.

Also, there are many plot types. We will see more in [Chapter 12](#), but another simple plot type is a **bar** chart.


For example, the following script creates two separate Figure Windows. First, it clears the Figure Window. Then, it creates an *x* vector and two different *y* vectors (*y1* and *y2*). In the first Figure Window, it plots the *y1* values using a bar chart. In the second Figure Window, it plots the *y1* values as black lines, puts **hold on** so that the next graph will be superimposed, and plots the *y2* values as black circles. It also puts a legend on this graph and uses a grid. Labels and titles are omitted in this case, as it is generic data.

plot2figs.m

```
% This creates 2 different plots, in 2 different
% Figure Windows, to demonstrate some plot features

clf
x = 1:5; % Not necessary
y1 = [2 11 6 9 3];
y2 = [4 5 8 6 2];
% Put a bar chart in Figure 1
figure(1)
bar(x,y1)
% Put plots using different y values on one plot
% with a legend
figure(2)
plot(x,y1,'k')
hold on
plot(x,y2,'ko')
grid on
legend('y1','y2')
```

Running this script will produce two separate Figure Windows. If there are no other active Figure Windows, the first, which is the bar chart, will be in the one titled “Figure 1” in MATLAB. The second will be in “Figure 2”. See [Figure 3.4](#) for both plots.

**FIGURE 3.4**

(A) Bar chart produced by script. (B) Plot produced by script, with a grid and legend.

Note that the first and last points are on the axes, which makes them difficult to see. That is why the `axis` function is used frequently, as it creates space around the points so that they are all visible.

PRACTICE 3.5


Modify the `plot2figs` script using the `axis` function so that all points are easily seen.

The ability to pass a vector to a function and have the function evaluate every element of the vector can be very useful in creating plots. For example, the following script graphically displays the difference between the `sin` and `cos` functions:

`sinnccos.m`

```
% This script plots sin(x) and cos(x) in the same Figure Window
% for values of x ranging from 0 to 2*pi

clf
x = 0: 2*pi/40: 2*pi;
y = sin(x);
plot(x,y,'ro')
hold on
y = cos(x);
plot(x,y,'b+')
legend('sin', 'cos')
xlabel('x')
ylabel('sin(x) or cos(x)')
title('sin and cos on one graph')
```

**FIGURE 3.5**

Plot of **sin** and **cos** in one Figure Window with a legend.

The script creates an x vector; iterating through all of the values from 0 to 2π in steps of $2\pi/40$ gives enough points to get a good graph. It then finds the sine of each x value and plots these. The command **hold on** freezes this in the Figure Window, so the next plot will be superimposed. Next, it finds the cosine of each x value and plots these points. The **legend** function creates a legend; the first character vector is paired with the first plot, and the second character vector with the second plot. Running this script produces the plot seen in Figure 3.5.

Beginning with Version R2014b, when **hold on** is used, MATLAB uses a sequence of colors for the plots, rather than using the default color for each. Of course, colors can also be specified as was done in this script.

Note that instead of using **hold on**, both functions could have been plotted using one call to the **plot** function:

```
plot(x, sin(x), x, cos(x))
```

PRACTICE 3.6

Write a script that plots **exp(x)** and **log(x)** for values of x ranging from 0 to 3.5.

3.6 INTRODUCTION TO FILE INPUT/OUTPUT (LOAD AND SAVE)

In many cases, input to a script will come from a data file that has been created by another source. Also, it is useful to be able to store output in an external file that can be manipulated and/or printed later. In this section, the simplest

methods used to read from an external data file and also to write to an external data file will be demonstrated.

There are basically three different operations, or *modes* on files. Files can be:

- read from
- written to
- appended to

Writing to a file means writing to a file from the beginning. *Appending to a file* is also writing, but starting at the end of the file rather than the beginning. In other words, appending to a file means adding to what was already there.

There are many different file types, which use different filename extensions. For now, we will keep it simple and just work with .dat or .txt files when working with data, or text, files. There are several methods for reading from files and writing to files; we will, for now, use the **load** function to read and the **save** function to write to files. More file types and functions for manipulating them will be discussed in [Chapter 9](#).

3.6.1 Writing Data to a File

The **save** command can be used to write data from a matrix to a data file, or to append to a data file. The format is:

```
save filename matrixvariablename -ascii
```

The “-ascii” qualifier is used when creating a text or data file. For example, the following creates a matrix and then saves the values from the matrix variable to a data file called *testfile.dat*:

```
>> mymat = rand(2,3)
mymat =
    0.4565    0.8214    0.6154
    0.0185    0.4447    0.7919
>> save testfile.dat mymat -ascii
```

This creates a file called “testfile.dat” that stores the numbers:

```
0.4565    0.8214    0.6154
0.0185    0.4447    0.7919
```

The **type** command can be used to display the contents of the file; note that scientific notation is used:

```
>> type testfile.dat

4.5646767e-001    8.2140716e-001    6.1543235e-001
1.8503643e-002    4.4470336e-001    7.9193704e-001
```


Note that if the file already exists, the **save** command will overwrite the file; **save** always writes from the beginning of a file.

3.6.2 Appending Data to a Data File

Once a text file exists, data can be appended to it. The format is the same as the preceding, with the addition of the qualifier “-append”. For example, the following creates a new random matrix and appends it to the file that was just created:

```
>> mat2 = rand(3,3)
mymat =
    0.9218    0.4057    0.4103
    0.7382    0.9355    0.8936
    0.1763    0.9169    0.0579
>> save testfile.dat mat2 -ascii -append
```

This results in the file “testfile.dat” containing the following:

```
    0.4565    0.8214    0.6154
    0.0185    0.4447    0.7919
    0.9218    0.4057    0.4103
    0.7382    0.9355    0.8936
    0.1763    0.9169    0.0579
```

Note that although technically any size matrix could be appended to this data file, to be able to read it back into a matrix later there would have to be the same number of values on every row (or, in other words, the same number of columns).

PRACTICE 3.7

Prompt the user for the number of rows and columns of a matrix, create a matrix with that many rows and columns of random integers, and write it to a file.

3.6.3 Reading From a File

Reading from a file is accomplished using **load**. Once a file has been created (as in the preceding), it can be read into a matrix variable. If the file is a data file, the **load** command will read from the file “filename.ext” (e.g., the extension might be .dat) and create a matrix with the same name as the file. For example, if the data file “testfile.dat” had been created as shown in the previous section, this would read from it and store the result in a matrix variable called *testfile*:

```
>> clear
>> load testfile.dat
>> who
Your variables are:
testfile
>> testfile
testfile =
    0.4565    0.8214    0.6154
    0.0185    0.4447    0.7919
    0.9218    0.4057    0.4103
    0.7382    0.9355    0.8936
    0.1763    0.9169    0.0579
```

The **load** command works only if there are the same number of values in each line, so that the data can be stored in a matrix, and the **save** command only writes from a matrix to a file. If this is not the case, lower-level file I/O functions must be used; these will be discussed in [Chapter 9](#).

3.6.3.1 Example: Load From a File and Plot the Data

As an example, a file called “timetemp.dat” stores two lines of data. The first line is the times of day, and the second line is the recorded temperature at each of those times. The first value of 0 for the time represents midnight. For example, the contents of the file might be:

```
0      3      6      9      12      15      18      21
55.5  52.4  52.6  55.7  75.6  77.7  70.3  66.6
```

The following script loads the data from the file into a matrix called *timetemp*. It then separates the matrix into vectors for the time and temperature, and then plots the data using black star (*) symbols.


```
timetempprob.m

% This reads time and temperature data for an afternoon
% from a file and plots the data

load timetemp.dat

% The times are in the first row, temps in the second row
time = timetemp(1,:);
temp = timetemp(2,:);

% Plot the data and label the plot
plot(time,temp,'k*')
xlabel('Time')
ylabel('Temperature')
title('Temperatures one afternoon')
```

**FIGURE 3.6**

Plot of temperature data from a file.

Running the script produces the plot seen in [Figure 3.6](#).

Note that it is difficult to see the point at time 0 as it falls on the y -axis. The `axis` function could be used to change the axes from the defaults shown here.

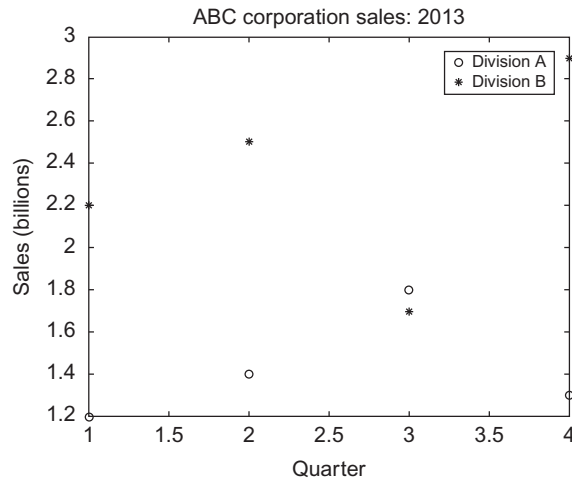
To create the data file, the Editor in MATLAB can be used; it is not necessary to create a matrix and **save** it to a file. Instead, just enter the numbers in a new script file, and Save As *timetemp.dat*, making sure that the Current Folder is set.

PRACTICE 3.8

The sales (in billions) for two separate divisions of the ABC Corporation for each of the four quarters of 2013 are stored in a file called "salesfigs.dat":

```
1.2 1.4 1.8 1.3
2.2 2.5 1.7 2.9
```

- First, create this file (just type the numbers in the Editor, and Save As "salesfigs.dat").
- Then, write a script that will
 - load the data from the file into a matrix
 - separate this matrix into 2 vectors.
 - create the plot seen in [Figure 3.7](#) (which uses black circles and stars as the plot symbols).

**FIGURE 3.7**

Plot of sales data from file.

QUICK QUESTION!

Sometimes files are not in the format that is desired. For example, a file “expresults.dat” has been created that has some experimental results, but the order of the values is reversed in the file:

```
4  53.4
3  44.3
2  50.0
1  55.5
```

How could we create a new file that reverses the order?

Answer: We can **load** from this file into a matrix, use the **flipud** function to “flip” the matrix up to down, and then **save** this matrix to a new file:

```
>> load expresults.dat
>> expresults
expresults =
    4.0000    53.4000
    3.0000    44.3000
    2.0000    50.0000
    1.0000    55.5000
>> correctorder = flipud(expresults)
correctorder =
    1.0000    55.5000
    2.0000    50.0000
    3.0000    44.3000
    4.0000    53.4000
>> save neworder.dat correctorder - ascii
```

3.7 USER-DEFINED FUNCTIONS THAT RETURN A SINGLE VALUE

We have already seen the use of many functions in MATLAB. We have used many built-in functions such as **sin**, **fix**, **abs**, and **double**. In this section, *user-defined functions* will be introduced. These are functions that the programmer defines, and then uses, in either the Command Window or in a script.

There are several different types of functions. For now, we will concentrate on the kind of function that calculates and returns a single result. Other types of functions will be introduced in [Chapter 6](#).

First, let us review some of what we already know about functions, including the use of built-in functions. Although by now the use of these functions is straightforward, explanations will be given in some detail here in order to compare to and contrast with the use of user-defined functions.

The **length** function is an example of a built-in function that calculates a single value; it returns the length of a vector. As an example,

```
length(vec)
```

is an expression that represents the number of elements in the vector *vec*. This expression could be used in the Command Window or in a script. Typically, the value returned from this expression might be assigned to a variable:

```
>> vec = 1:3:10;  
>> lv = length(vec)  
lv =  
4
```

Alternatively, the length of the vector could be printed:

```
>> fprintf('The length of the vector is %d\n', length(vec))  
The length of the vector is 4
```

The **function call** to the **length** function consists of the name of the function, followed by the **argument** in parentheses. The function receives as input the argument and returns a result. What happens when the call to the function is encountered is that **control** is passed to the function itself (in other words, the function begins executing). The argument(s) are also passed to the function.

The function executes its statements and does whatever is necessary (the actual contents of the built-in functions are not generally known or seen by the user) to determine the number of elements in the vector. As the function is calculating a single value, this result is then **returned** and it becomes the value of the expression. Control is also passed back to the expression that called it in the first place, which then continues (e.g., in the first example the value would then be assigned to the variable *lv* and in the second example the value was printed).

3.7.1 Function Definitions

There are different ways to organize scripts and functions, but, for now, every function that we write will be stored in a separate file. Like scripts, function files have an extension of *.m*. Although to enter function definitions in the Editor it is possible to choose the New down arrow and then Function, it will be easier

for now to type in the function by choosing New Script (this ignores the defaults that are provided when you choose Function).

A function in MATLAB that returns a single result consists of the following.

- The *function header* (the first line), comprised of:
 - the reserved word **function**
 - the name of the *output argument* followed by the assignment operator (=), as the function *returns* a result
 - the name of the function (*important*—This should be the same as the name of the file in which this function is stored to avoid confusion)
 - the *input arguments* in parentheses, which correspond to the arguments that are passed to the function in the function call
- A comment that describes what the function does (this is printed when **help** is used)
- The *body* of the function, which includes all statements and eventually must put a value in the output argument
- **end** at the end of the function (note that this is not necessary in many cases in current versions of MATLAB, but it is considered good style anyway)

The general form of a *function definition* for a function that calculates and returns one value looks like this:

```
functionname.m
```

```
function outputargument = functionname(input arguments)
% Comment describing the function

Statements here; these must include putting a value in the output
argument

end % of the function
```

For example, the following is a function called *calcarearea* that calculates and returns the area of a circle; it is stored in a file called *calcarearea.m*.

```
calcarearea.m
```

```
function area = calcarea(rad)
% calcarea calculates the area of a circle
% Format of call: calcarea(radius)
% Returns the area

area = pi * rad * rad;
end
```

A radius of a circle is passed to the function to the input argument *rad*; the function calculates the area of this circle and stores it in the output argument *area*.

In the function header, we have the reserved word **function**, then the output argument *area* followed by the assignment operator `=`, then the name of the function (the same as the name of the file), and then the input argument *rad*, which is the radius. As there is an output argument in the function header, somewhere in the body of the function we must put a value in this output argument. This is how a value is returned from the function. In this case, the function is simple and all we have to do is assign to the output argument *area* the value of the built-in constant **pi** multiplied by the square of the input argument *rad*.

The function can be displayed in the Command Window using the **type** command.

```
>> type calcarearea

function area = calcarearea(rad)
% calcarearea calculates the area of a circle
% Format of call: calcarearea(radius)
% Returns the area

area = pi * rad * rad;
end
```

Note that many of the functions in MATLAB are implemented as functions that are stored in files with an extension of `.m`; these can also be displayed using **type**.

3.7.2 Calling a Function

The following is an example of a call to this function in which the value returned is stored in the default variable *ans*:

```
>> calcarearea(4)
ans =
    50.2655
```

Technically, calling the function is done with the name of the file in which the function resides. To avoid confusion, it is easiest to give the function the same name as the file name, so that is how it will be presented in this book. In this example, the function name is *calcarearea* and the name of the file is *calcarearea.m*. The result returned from this function can also be stored in a variable in an assignment statement; the name could be the same as the name of the output argument in the function itself, but that is not necessary. So, for example, either of these assignments would be fine:

```
>> area = calcarearea(5)
area =
    78.5398
```

```
>> myarea = calcarearea(6)
myarea =
    113.0973
```

The output could also be suppressed when calling the function:

```
>> mya = calcarearea(5.2);
```

The value returned from the *calcarearea* function could also be printed using either **disp** or **fprintf**:

```
>> disp(calcarearea(4))
    50.2655
>> fprintf('The area is %.1f\n', calcarearea(4))
The area is 50.3
```

Note that the printing is not done in the function itself; rather, the function returns the area and then an output statement can print or display it.

QUICK QUESTION!

Could we pass a vector of radii to the *calcarearea* function?

Answer: This function was written assuming that the argument was a scalar, so calling it with a vector instead would produce an error message:

```
>> calcarearea(1:3)
Error using *
    Inner matrix dimensions must agree.

Error in calcarearea (line 6)
    area = pi * rad * rad;
```

This is because the `*` was used for multiplication in the function, but `.*` must be used when multiplying vectors term by term. Changing this in the function would allow either scalars or vectors to be passed to this function:

calcarearea.m

```
function area = calcarearea(rad)
% calcarearea returns the area of a circle
% The input argument can be a vector of radii
% Format: calcarearea(radiiVector)

area = pi * rad .* rad;
end
```

```
>> calcarearea(1:3)
ans =
    3.1416   12.5664   28.2743

>> calcarearea(4)
ans =
    50.2655
```

Note that the `.*` operator is only necessary when multiplying the radius vector by itself. Multiplying by **pi** is scalar multiplication, so the `.*` operator is not needed there. We could have also used:

```
area = pi * rad . ^ 2;
```


Using **help** with either of these functions displays the contiguous block of comments under the function header (the block comment). It is useful to put the format of the call to the function in this block comment:

```
>> help calcareea
    calcareea calculates the area of a circle
    Format of call: calcareea(radius)
    Returns the area
```

The suggested corrections for invalid filenames in the Command Window work for user-defined files as of Version R2014b.

```
>> clacarea(3)
Undefined function or variable 'clacarea'.
Did you mean:
>> calcareea(3)
```

Many organizations have standards regarding what information should be included in the block comment in a function. These can include:

- Name of the function
- Description of what the function does
- Format of the function call
- Description of input arguments
- Description of output argument
- Description of variables used in function
- Programmer name and date written
- Information on revisions

Although this is excellent programming style, for the most part in this book these will be omitted simply to save space. Also, documentation in MATLAB suggests that the name of the function should be in all uppercase letters in the beginning of the block comment. However, this can be somewhat misleading in that MATLAB is case-sensitive and typically lowercase letters are used for the actual function name.

3.7.3 Calling a User-Defined Function From a Script

Now, we will modify our script that prompts the user for the radius and calculates the area of a circle to call our function *calcareea* to calculate the area of the circle rather than doing this in the script.

```

circleCallFn.m

% This script calculates the area of a circle
% It prompts the user for the radius
radius = input('Please enter the radius: ');
% It then calls our function to calculate the
% area and then prints the result
area = calcarea(radius);
fprintf('For a circle with a radius of %.2f, ', radius)
fprintf(' the area is %.2f\n', area)

```

Running this will produce the following:

```

>> circleCallFn
Please enter the radius: 5
For a circle with a radius of 5.00, the area is 78.54

```

3.7.3.1 Simple Programs

In this book, a script that calls function(s) is what we will call a MATLAB program. In the previous example, the program consisted of the script *circleCallFn* and the function it calls, *calcarea*. A simple program, consisting of a script that calls a function to calculate and return a value, looks like the format shown in Figure 3.8.

It is also possible for a function to call another (whether built-in or user-defined).

3.7.4 Passing Multiple Arguments

In many cases, it is necessary to pass more than one argument to a function. For example, the volume of a cone is given by

$$V = \frac{1}{3}\pi r^2 h$$




FIGURE 3.8

General form of a simple program.

where r is the radius of the circular base and h is the height of the cone. Therefore, a function that calculates the volume of a cone needs both the radius and the height:

```
conevol.m

function outarg = conevol(radius, height)
% conevol calculates the volume of a cone
% Format of call: conevol(radius, height)
% Returns the volume

outarg = (pi/3) * radius . ^ 2 .* height;
end
```

As the function has two input arguments in the function header, two values must be passed to the function when it is called. The order makes a difference. The first value that is passed to the function is stored in the first input argument (in this case, *radius*) and the second argument in the function call is passed to the second input argument in the function header.

This is very important: the arguments in the function call must correspond one-to-one with the input arguments in the function header.

Here is an example of calling this function. The result returned from the function is simply stored in the default variable *ans*.

```
>> conevol(4,6.1)
ans =
    102.2065
```

In the next example, the result is instead printed with a format of two decimal places.

```
>> fprintf('The cone volume is %.2f\n', conevol(3, 5.5))
The cone volume is 51.84
```

Note that by using the array exponentiation and multiplication operators, it would be possible to pass arrays for the input arguments, as long as the dimensions are the same.

QUICK QUESTION!

Nothing is technically wrong with the following function, but what if it does not make sense?

Answer: Why pass the third argument if it is not used?

```
fun.m

function out = fun(a,b,c)
out = a*b;
end
```

PRACTICE 3.9

Write a script that will prompt the user for the radius and height, call the function *conevol* to calculate the cone volume, and print the result in a nice sentence format. So, the program will consist of a script and the *conevol* function that it calls.

PRACTICE 3.10

For a project, we need some material to form a rectangle. Write a function *calcrectarea* that will receive the length and width of a rectangle in inches as input arguments and will return the area of the rectangle. For example, the function could be called as shown, in which the result is stored in a variable and then the amount of material required is printed, rounded up to the nearest square inch.

```
>> ra = calcrectarea(3.1, 4.4)
ra =
    13.6400

>> fprintf('We need %d sq in.\n', ceil(ra))
We need 14 sq in.
```

3.7.5 Functions With Local Variables

The functions discussed thus far have been very simple. However, in many cases the calculations in a function are more complicated and may require the use of extra variables within the function; these are called *local variables*.

For example, a closed cylinder is being constructed of a material that costs a certain dollar amount per square foot. We will write a function that will calculate and return the cost of the material, rounded up to the nearest square foot, for a cylinder with a given radius and a given height. The total surface area for the closed cylinder is

$$SA = 2\pi rh + 2\pi r^2$$

For a cylinder with a radius of 32 inches, height of 73 inches, and cost per square foot of the material of \$4.50, the calculation would be given by the following algorithm.

- Calculate the surface area $SA = 2*\pi*32*73 + 2*\pi*32*32$ inches squared.
- Convert the SA from square inches to square feet = $SA/144$.
- Calculate the total cost = SA in square feet * cost per square foot.

The function includes local variables to store the intermediate results.

```
cylcost.m

function outcost = cylcost(radius, height, cost)
% cylcost calculates the cost of constructing a closed
% cylinder
% Format of call: cylcost(radius, height, cost)
% Returns the total cost

% The radius and height are in inches
% The cost is per square foot

% Calculate surface area in square inches
surf_area = 2 * pi * radius .* height + 2 * pi * radius . ^ 2;

% Convert surface area in square feet and round up
surf_areasf = ceil(surf_area/144);

% Calculate cost
outcost = surf_areasf .* cost;
end
```

The following shows examples of calling the function:

```
>> cylcost(32, 73, 4.50)
ans =
    661.5000

>> fprintf('The cost would be $%.2f\n', cylcost(32, 73, 4.50))
The cost would be $661.50
```

3.7.6 Introduction to Scope

It is important to understand the *scope of variables*, which is where they are valid. More will be described in [Chapter 6](#), but, basically, variables used in a script are also known in the Command Window and vice versa. All variables used in a function, however, are local to that function. Both the Command Window and scripts use a common workspace, the *base workspace*. Functions, however, have their own workspaces. This means that when a script is executed, the variables can subsequently be seen in the Workspace Window and can be used from the Command Window. This is not the case with functions, however.

3.8 LOCAL FUNCTIONS IN SCRIPTS

In [Section 3.7.3](#), a program was described as a script that calls a function; both the script and the function were stored in separate code files. [Figure 3.8](#) illustrated this organization. However, as of R2016b, it is now possible to include functions within a script file, at the end of the script file. Functions that are

contained within a script code file are *local functions* to that script and can only be called from within the script code file. They are sometimes called *subfunctions*. All local functions, like other user-defined functions, have their own workspaces. Even though a local function is contained in a script, the script uses the base workspace, whereas local functions use their own workspaces.

The following script code file includes a local function:

```
testlocfn.m

x = 33;
y = 11;

a = locfn(x);
fprintf('a is %d\n', a)
fprintf('x is %d\n', x)

function out = locfn(in)
x = in + 5;
out = x;
end
```

The script will create the variables x , y , and a in the base workspace. While the function is executing, it will have in , x , and out in its workspace, but that workspace only exists while the function is executing. Therefore, the value of x that is printed from the script is the value assigned in the script in the base workspace.

```
>> testlocfn
a is 38
x is 33
```

An advantage of having a local function is that it cuts down on the number of code files. If a function is only going to be called from a script (and it is not desired to call it from the Command Window or from another function), then it is easier to include it in the same code file as the script. Another use for a local function is to have an easy way of testing the function when the function is first being written; later when finished, the function could be placed in its own code file. The disadvantage of a local function is that it can only be called from within the script code file. This means that it can be called from within the script, or from another local function within the script. The order of the local functions does not matter if there are multiple local functions, but they must appear after the script code.

3.9 COMMANDS AND FUNCTIONS

Some of the commands that we have used (e.g., **format**, **type**, **save**, and **load**) are just shortcuts for function calls. If all of the arguments to be passed to a

function are strings or character vectors, and the function does not return any values, it can be used as a command. For example, the following produce the same results:

```
>> type script1

radius = 5
area = pi * (radius ^2)

>> type('script1')

radius = 5
area = pi * (radius ^2)
```

Using **load** as a command creates a variable with the same name as the file. If a different variable name is desired, it is easiest to use the functional form of **load**. For example,

```
>> type pointcoords.dat

3.3    1.2
4      5.3

>> points = load('pointcoords.dat')
points =
    3.3000    1.2000
    4.0000    5.3000
```

This stores the result in a variable *points* rather than *pointcoords*.

■ Explore Other Interesting Features

Note that this chapter serves as an introduction to several topics, most of which will be covered in more detail in future chapters. Before getting to those chapters, the following are some things you may wish to explore.

- The **help** command can be used to see short explanations of built-in functions. At the end of this, a doc page link is also listed. These documentation pages frequently have much more information and useful examples. They can also be reached by typing “doc fname” where fname is the name of the function.
- Look at formatSpec on the doc page on the **sprintf** function for more ways in which expressions can be formatted, e.g., padding numbers with zeros and printing the sign of a number.
- Use the Search Documentation to find the conversion characters used to print other types, such as unsigned integers and exponential notation. ■

SUMMARY

COMMON PITFALLS

- Spelling a variable name different ways in different places in a script or function.
- Forgetting to add the second 's' argument to the **input** function when character input is desired.
- Not using the correct conversion character when printing.
- Confusing **fprintf** and **disp**. Remember that only **fprintf** can format.

PROGRAMMING STYLE GUIDELINES

- Especially for longer scripts and functions, start by writing an algorithm.
- Use comments to document scripts and functions, as follows:
 - a block of contiguous comments at the top to describe a script
 - a block of contiguous comments under the function header for functions
 - comments throughout any code file (script or function) to describe each section
- Make sure that the "H1" comment line has useful information.
- Use your organization's standard style guidelines for block comments.
- Use mnemonic identifier names (names that make sense, e.g., *radius* instead of *xyz*) for variable names and for file names.
- Make all output easy to read and informative.
- Put a newline character at the end of every format specifier printed by **fprintf** so that the next output or the prompt appears on the line below.
- Put informative labels on the *x*- and *y*-axes, and a title on all plots.
- Keep functions short—typically no longer than one page in length.
- Suppress the output from all assignment statements in functions and scripts.
- Functions that return a value do not normally print the value; it should simply be returned by the function.
- Use the array operators **.***, **./**, **.**, and **.^** in functions so that the input arguments can be arrays and not just scalars.

MATLAB Reserved Words

function	end
----------	-----

MATLAB Functions and Commands

type	plot	clf	bar
input	xlabel	figure	load
disp	ylabel	hold	save
fprintf	title	legend	
display	axis	grid	

MATLAB Operators

comment %	comment block %{, %}
-----------	----------------------

Exercises

- Using the top-down design approach, write an algorithm for making a sandwich.
- Write a simple script that will calculate the volume of a hollow sphere,

$$\frac{4\pi}{3}(r_0^3 - r_i^3)$$

where r_i is the inner radius and r_0 is the outer radius. Assign a value to a variable for the inner radius, and also assign a value to another variable for the outer radius. Then, using these variables, assign the volume to a third variable. Include comments in the script. Use **help** to view the block comment in your script.

- Write a statement that prompts the user for his/her favorite number.
- Write a statement that prompts the user for his/her name.
- Write an **input** statement that will prompt the user for a real number and store it in a variable. Then, use the **fprintf** function to print the value of this variable using 2 decimal places.
- Experiment, in the Command Window, with using the **fprintf** function for real numbers. Make a note of what happens for each. Use **fprintf** to print the real number 12345.6789
 - without specifying any field width
 - in a field width of 10 with 4 decimal places
 - in a field width of 10 with 2 decimal places
 - in a field width of 6 with 4 decimal places
 - in a field width of 2 with 4 decimal places

7. Experiment, in the Command Window, with using the **fprintf** function for integers. Make a note of what happens for each. Use **fprintf** to print the integer 12345
 - without specifying any field width
 - in a field width of 5
 - in a field width of 8
 - in a field width of 3
8. When would you use **disp** instead of **fprintf**? When would you use **fprintf** instead of **disp**?
9. Write a script called *echostring* that will prompt the user for a character vector and will echo print it in quotes:

```
>> echostring
Enter your string: hi there
Your string was: 'hi there'
```

10. Experiment with the input function to determine how the user could enter a matrix.
11. Trace the following code and determine what will be printed, and then enter it in a script and execute it to verify your results.

```
num = 12.3;
disp('Hello!')
fprintf('Hmm how many ')
fprintf('lines\n\nwill print?\n')
fprintf('%6.2f\n', num)
```

12. Write a script that will prompt the user for a character and will print it twice; once left-justified in a field width of 5, and again right-justified in a field width of 3.
13. A script “iotrace” has been written. Here’s what the desired output looks like:

```
>> iotrace
Please enter a number: 33
Please enter a character: x
Your number is 33.00
Your char is  x!
```

Fix this script so that it works as shown previously:

```
mynum = input('Please enter a number:\n ');
mychar = input('Please enter a character: ');
fprintf('Your number is %.2f, mynum)
fprintf('Your char is %c!\n', mychar)
```

14. A power company charges 6.6 cents per kilowatt hour (KWH) for providing electricity. Write a script “power_charge” that will prompt the user for the number of KWH used in a given month and will print the charge for the month in dollars, in the following format. (The conversion is 100 cents in one dollar.) Here are examples of using the script.

```
>> power_charge
How many KWH this month: 200
Your charge for 200 KWH will be $13.20.
>> help power_charge
Calculates and prints charge to customer for electricity
>>
```

15. Why do we always suppress all assignment statements in scripts?
16. For a bubble, the surface tension force in the downward direction is

$$F_d = 4\pi Tr$$

where T is the surface tension measured in force per unit length and r is the radius of the bubble. For water, the surface tension at 25°C is 72 dyn/cm. Write a script *surftens* that will prompt the user for the radius of a water bubble in centimeters, calculate F_d , and print it in a sentence (ignoring units for simplicity). Assume that the temperature of the water is 25°C, so use 72 for T .

17. Write a script that assigns values for the x coordinate and then y coordinate of a point, and then plots this using a green +.
18. Plot **sin(x)** for x values ranging from 0 to π (in separate Figure Windows):
 - using 10 points in this range
 - using 100 points in this range
19. Write a script that will:
 - prompt the user for a maximum x value
 - create an x vector from 1 to the user's maximum in steps of 0.1
 - create a y vector which is $\sin(x)$
 - plot the x and y vectors using blue *'s , using appropriate x and y labels and a title
20. When would it be important to use **legend** in a plot?
21. Generate a random integer n , create a vector of the integers 1 through n in steps of 2, square them, and plot the squares.
22. Atmospheric properties such as temperature, air density, and air pressure are important in aviation. Create a file that stores temperatures in degrees Kelvin at various altitudes. The altitudes are in the first column and the temperatures in the second. For example, it may look like this:

1000	288
2000	281
3000	269
23. Create a 3×6 matrix of random integers, each in the range from 50 to 100. Write this to a file called *randfile.dat*. Then, create a new matrix of random integers, but this time make it a 2×6 matrix of random integers, each in the range from 50 to 100. Append this matrix to the original file. Then, read the file in (which will be to a variable called *randfile*) just to make sure that worked!

24. Write a script that would create a 5×5 matrix of random integers. From this, create another matrix variable which is the “middle” 3×3 part of the original matrix. Write this new matrix to a file.
25. A part is being turned on a lathe. The diameter of the part is supposed to be 20,000mm. The diameter is measured every 10 minutes and the results are stored in a file called *partdiam.dat*. Create a data file to simulate this. The file will store the time in minutes and the diameter at each time. Plot the data.
26. List some differences between a script and a function.
27. Write a function that converts inches to feet. Note that 12 inches = 1 foot.
28. Write a *fives* function that will receive two arguments for the number of rows and columns and will return a matrix with that size of all fives.
29. A calorie is a unit of energy. The 15°C calorie is defined as 4.1855 Joules. The following script defines a couple of random calorie values, calls a function to calculate the equivalent number of Joules, and prints this. Write the function and store it in a separate code file.

```
cal1 = randi([10 2e3]);
inj1 = cal2J(cal1);
cal2 = rand*100;
inj2 = cal2J(cal2);
fprintf('%d calories is %.2fJ\n', cal1, inj1)
fprintf('%d calories is %.2fJ\n', cal2, inj2)
```

30. Write a function called *rotleft* that will receive one row vector as an argument (you may assume that it is a row vector with a length of at least two) and will return another vector, which consists of the input vector rotated to the left—e.g., all values shift over one element, and the first element is wrapped around to the end. For example,

```
>> rotleft([1 3 4])
ans =
     3     4     1
```

31. Write a function *isdivby4* that will receive an integer input argument and will return **logical 1** for **true** if the input argument is divisible by 4, or **logical false** if it is not.
32. Write a function *isint* that will receive a number input argument *innum* and will return 1 for **true** if this number is an integer, or 0 for **false** if not. Use the fact that *innum* should be equal to **int32(innum)** if it is an integer. Unfortunately, due to round-off errors, it should be noted that it is possible to get **logical 1** for **true** if the input argument is close to an integer. Therefore, the output may not be what you might expect, as shown here.

where A is the initial value at $t=0$, and τ is the time constant for the function. Write a script to study the effect of the time constant. To simplify the equation, set A equal to 1. Prompt the user for two different values for the time constant and for beginning and ending values for the range of a t vector. Then, calculate two different y vectors using the above equation and the two time constants, and graph both exponential functions on the same graph within the range the user specified. Use a function to calculate y . Make one plot red. Be sure to label the graph and both axes. What happens to the decay rate as the time constant gets larger?